

# EE 628

# Deep Learning

# Fall 2019

Lecture 6  
02/27/2019

Applied Scientist Sergul Aydore  
*Amazon Web Services*

# Overview

- Last lecture we covered
  - Backpropagation
  - Underfitting/Overfitting
- Today, we will cover
  - Optimization Algorithms
  - Convolutional layers

# Optimization Algorithms

- Optimization Algorithms are important for deep learning

# Optimization Algorithms

- Optimization Algorithms are important for deep learning
- The performance of the training algorithm directly affects the model's training efficiency

# Optimization Algorithms

- Optimization Algorithms are important for deep learning
- The performance of the training algorithm directly affects the model's training efficiency
- Understanding the principles of different optimization algorithms will help us tune the parameters in a targeted manner

# Optimization Algorithms

- Optimization Algorithms are important for deep learning
- The performance of the training algorithm directly affects the model's training efficiency
- Understanding the principles of different optimization algorithms will help us tune the parameters in a targeted manner
- Now, it is time to explore common deep learning algorithms in depth

# Optimization Algorithms

- Optimization Algorithms are important for deep learning
- The performance of the training algorithm directly affects the model's training efficiency
- Understanding the principles of different optimization algorithms will help us tune the parameters in a targeted manner
- Now, it is time to explore common deep learning algorithms in depth
- Almost all problems arising in deep learning are nonconvex

# Optimization Algorithms

- Optimization Algorithms are important for deep learning
- The performance of the training algorithm directly affects the model's training efficiency
- Understanding the principles of different optimization algorithms will help us tune the parameters in a targeted manner
- Now, it is time to explore common deep learning algorithms in depth
- Almost all problems arising in deep learning are nonconvex
- However, analysis of the algorithms in the context of convex problems can be very instructive.



# Optimization and Deep Learning

- We usually define a loss function first in deep learning

# Optimization and Deep Learning

- We usually define a loss function first in deep learning
- Then, we use an optimization algorithm in an attempt to minimize it

# Optimization and Deep Learning


- We usually define a loss function first in deep learning
- Then, we use an optimization algorithm in an attempt to minimize it
- By tradition most optimization algorithms are concerned with minimization

# Optimization and Deep Learning

- We usually define a loss function first in deep learning
- Then, we use an optimization algorithm in an attempt to minimize it
- By tradition most optimization algorithms are concerned with minimization
- The goal of optimization is to reduce training error but the goal is deep learning (statistical inference in general)

# Optimization and Deep Learning

- We usually define a loss function first in deep learning
- Then, we use an optimization algorithm in an attempt to minimize it
- By tradition most optimization algorithms are concerned with minimization
- The goal of optimization is to reduce training error but the goal is deep learning (statistical inference in general)




# Optimization Challenges in Deep Learning

- We are going to focus specifically on the performance of the optimization algorithm in minimizing the objective function



# Optimization Challenges in Deep Learning

- We are going to focus specifically on the performance of the optimization algorithm in minimizing the objective function
- **Local Minima:** The objective functions of deep learning models usually has many local minima.
  - Only some degree of noise might knock the parameter out of the local minimum.
  - This is one of the beneficial properties of SGD.



# Optimization Challenges in Deep Learning

- We are going to focus specifically on the performance of the optimization algorithm in minimizing the objective function
- **Local Minima:** The objective functions of deep learning models usually has many local minima.
  - Only some degree of noise might knock the parameter out of the local minimum.
  - This is one of the beneficial properties of SGD.
- **Saddle Points:** Any location where all gradients vanish but which is neither a global or a local minimum.
  - Optimization might stall at the point even though it is not a minimum.






# Optimization Challenges in Deep Learning


- We are going to focus specifically on the performance of the optimization algorithm in minimizing the objective function
- **Local Minima:** The objective functions of deep learning models usually has many local minima.

- Only some degree of noise might knock the parameter out of the local minimum.
- This is one of the beneficial properties of SGD.



- **Saddle Points:** Any location where all gradients vanish but which is neither a global or a local minimum.

- Optimization might stall at the point even though it is not a minimum.



- **Vanishing Gradients:** Probably, the most insidious problem

# Convexity

- It is much easier to design and test algorithms in the context of convexity.

# Convexity

- It is much easier to design and test algorithms in the context of convexity.
- Optimization problems in DL still exhibit some properties of convex problems near local minima




# Convexity

- It is much easier to design and test algorithms in the context of convexity.
- Optimization problems in DL still exhibit some properties of convex problems near local minima
- **Sets:** A set  $X$  in a vector space is convex if for any  $a, b \in X$  the line segment connecting  $a$  and  $b$  is also in  $X$ . Mathematically, for all  $\lambda \in [0, 1]$  we have

$$\lambda a + (1 - \lambda)b \in X \text{ whenever } a, b \in X$$

# Convexity

- It is much easier to design and test algorithms in the context of convexity.
- Optimization problems in DL still exhibit some properties of convex problems near local minima
- **Sets:** A set  $X$  in a vector space is convex if for any  $a, b \in X$  the line segment connecting  $a$  and  $b$  is also in  $X$ . Mathematically, for all  $\lambda \in [0, 1]$  we have



# Convexity

- Say  $X$  and  $Y$  are convex sets, then  $X \cap Y$  is ?

# Convexity

- Say  $X$  and  $Y$  are convex sets, then  $X \cap Y$  is ?
- Say  $X$  and  $Y$  are convex sets, then  $X \cup Y$  is ?

# Convexity

- Say  $X$  and  $Y$  are convex sets, then  $X \cap Y$  is ?
- Say  $X$  and  $Y$  are convex sets, then  $X \cup Y$  is ?
- Typically, the problems in DL are defined on convex domains
  - For instance  $\mathbb{R}^d$  is a convex set
  - In some cases we work with variables of bounded length, such as balls of radius  $r$

$$\{\mathbf{x} | \mathbf{x} \in \mathbb{R}^d \text{ and } \|\mathbf{x}\|_2 \leq r\}$$




# Convexity

- Say  $X$  and  $Y$  are convex sets, then  $X \cap Y$  is ?
- Say  $X$  and  $Y$  are convex sets, then  $X \cup Y$  is ?
- Typically, the problems in DL are defined on convex domains
  - For instance  $\mathbb{R}^d$  is a convex set
  - In some cases we work with variables of bounded length, such as balls of radius  $r$

$$\{\mathbf{x} | \mathbf{x} \in \mathbb{R}^d \text{ and } \|\mathbf{x}\|_2 \leq r\}$$

- **Functions:** Given a convex set  $X$ , a function defined on it  $f: X \rightarrow \mathbb{R}$  is convex if for all  $x, x' \in X$  and for all  $\lambda \in [0, 1]$  we have

$$\lambda f(x) + (1 - \lambda)f(x') \geq f(\lambda x + (1 - \lambda)x')$$



# Convexity

- Jensen's Inequality: It amounts to generalization of the definition of convexity.

$$\sum_i \alpha_i f(x_i) \geq f\left(\sum_i \alpha_i x_i\right) \text{ and } \mathbb{E}_x[f(x)] \geq f(\mathbb{E}_x[x])$$

# Convexity

- Jensen's Inequality: It amounts to generalization of the definition of convexity. 
$$\sum_i \alpha_i f(x_i) \geq f\left(\sum_i \alpha_i x_i\right) \text{ and } \mathbb{E}_x[f(x)] \geq f(\mathbb{E}_x[x])$$
- In other words, the expectation of a convex function is larger than the convex function of an expectation.

# Convexity


- Jensen's Inequality: It amounts to generalization of the definition of convexity. 
$$\sum_i \alpha_i f(x_i) \geq f\left(\sum_i \alpha_i x_i\right) \text{ and } \mathbb{E}_x[f(x)] \geq f(\mathbb{E}_x[x])$$
- In other words, the expectation of a convex function is larger than the convex function of an expectation.
- Can you prove this?

# Gradient Descent

- Let's start with an example in one dimension to explain why the gradient descent algorithm may reduce the value of the objective function.
  - Prove
  - Hint: Use Taylor's series expansion around  $x + \epsilon$  and then replace  $\epsilon$  with  $-\eta f'(x)$ .



# Gradient Descent

- Let's start with an example in one dimension to explain why the gradient descent algorithm may reduce the value of the objective function.
  - Prove
  - Hint: Use Taylor's series expansion around  $x + \epsilon$  and then replace  $\epsilon$  with  $-\eta f'(x)$ .



# Learning Rate in Gradient Descent

- Which one has large learning rate and which one has small?



# Multivariate Gradient Descent

- Let us consider a situation where  $x \in R^d$



# Multivariate Gradient Descent

- Let us consider a situation where  $x \in R^d$
- The objective function  $f: R^d \rightarrow R$  maps vectors into scalars

# Multivariate Gradient Descent

- Let us consider a situation where  $x \in R^d$
- The objective function  $f: R^d \rightarrow R$  maps vectors into scalars
- Correspondingly, its gradient is a vector of consisting  $d$  partial derivatives

$$\nabla f(\mathbf{x}) = \left[ \frac{\partial f(\mathbf{x})}{\partial x_1}, \frac{\partial f(\mathbf{x})}{\partial x_2}, \dots, \frac{\partial f(\mathbf{x})}{\partial x_d} \right]^\top.$$

# Multivariate Gradient Descent

- Let us consider a situation where  $x \in R^d$
- The objective function  $f: R^d \rightarrow R$  maps vectors into scalars
- Correspondingly, its gradient is a vector of consisting  $d$  partial derivatives

$$\nabla f(\mathbf{x}) = \left[ \frac{\partial f(\mathbf{x})}{\partial x_1}, \frac{\partial f(\mathbf{x})}{\partial x_2}, \dots, \frac{\partial f(\mathbf{x})}{\partial x_d} \right]^\top.$$

- As before in the univariate case, we can use the corresponding Taylor approximation for multivariate functions

$$f(\mathbf{x} + \epsilon) = f(\mathbf{x}) + \epsilon^\top \nabla f(\mathbf{x}) + O(\|\epsilon\|^2).$$

# Multivariate Gradient Descent

- Let us consider a situation where  $x \in R^d$
- The objective function  $f: R^d \rightarrow R$  maps vectors into scalars
- Correspondingly, its gradient is a vector of consisting  $d$  partial derivatives

$$\nabla f(\mathbf{x}) = \left[ \frac{\partial f(\mathbf{x})}{\partial x_1}, \frac{\partial f(\mathbf{x})}{\partial x_2}, \dots, \frac{\partial f(\mathbf{x})}{\partial x_d} \right]^\top.$$

- As before in the univariate case, we can use the corresponding Taylor approximation for multivariate functions

$$f(\mathbf{x} + \epsilon) = f(\mathbf{x}) + \epsilon^\top \nabla f(\mathbf{x}) + O(\|\epsilon\|^2).$$

- Choosing a positive learning rate yields the gradient descent algorithm

$$\mathbf{x} \leftarrow \mathbf{x} - \eta \nabla f(\mathbf{x}).$$

# Multivariate Gradient Descent

- Let  $f(\mathbf{x}) = x_1^2 + 2x_2^2$  be the objective function

# Multivariate Gradient Descent


- Let  $f(\mathbf{x}) = x_1^2 + 2x_2^2$  be the objective function
- What is the gradient?

# Multivariate Gradient Descent

- Let  $f(\mathbf{x}) = x_1^2 + 2x_2^2$  be the objective function
- What is the gradient?
- If the initial position is  $[-5, -2]$ , what is the next position?

# Multivariate Gradient Descent

- Let  $f(\mathbf{x}) = x_1^2 + 2x_2^2$  be the objective function
- What is the gradient?
- If the initial position is  $[-5, -2]$ , what is the next position?





# Adaptive Methods

- Getting the learning rate  $\eta$  just right is tricky

# Adaptive Methods

- Getting the learning rate  $\eta$  just right is tricky
- What if we determine  $\eta$  automatically?

# Adaptive Methods

- Getting the learning rate  $\eta$  just right is tricky
- What if we determine  $\eta$  automatically?
- Second order methods that also look at the value of *curvature* can help.

# Adaptive Methods

- Getting the learning rate  $\eta$  just right is tricky
- What if we determine  $\eta$  automatically?
- Second order methods that also look at the value of *curvature* can help.
- They cannot be applied directly to DL due to computational cost

# Adaptive Methods

- Getting the learning rate  $\eta$  just right is tricky
- What if we determine  $\eta$  automatically?
- Second order methods that also look at the value of *curvature* can help.
- They cannot be applied directly to DL due to computational cost
- But they provide useful intuition into how to design advanced optimization algorithms

# Newton's method

- Reviewing the Taylor expansion of  $f$  there is no need to stop after the first term. In fact, we can write it as

$$f(\mathbf{x} + \epsilon) = f(\mathbf{x}) + \epsilon^\top \nabla f(\mathbf{x}) + \frac{1}{2} \epsilon^\top \nabla \nabla^\top f(\mathbf{x}) \epsilon + O(\|\epsilon\|^3)$$

# Newton's method

- Reviewing the Taylor expansion of  $f$  there is no need to stop after the first term. In fact, we can write it as

$$f(\mathbf{x} + \epsilon) = f(\mathbf{x}) + \epsilon^\top \nabla f(\mathbf{x}) + \frac{1}{2} \epsilon^\top \nabla \nabla^\top f(\mathbf{x}) \epsilon + O(\|\epsilon\|^3)$$

- We define  $H_f := \Delta \Delta^\top f(\mathbf{x})$  to be the Hessian of  $f$ 
  - Easy to compute for small  $d$  and simple problems
  - For DL,  $H_f$  maybe too large and expensive to compute via backprop

# Newton's method

- Reviewing the Taylor expansion of  $f$  there is no need to stop after the first term. In fact, we can write it as

$$f(\mathbf{x} + \epsilon) = f(\mathbf{x}) + \epsilon^\top \nabla f(\mathbf{x}) + \frac{1}{2} \epsilon^\top \nabla \nabla^\top f(\mathbf{x}) \epsilon + O(\|\epsilon\|^3)$$

- We define  $H_f := \Delta \Delta^\top f(\mathbf{x})$  to be the Hessian of  $f$ 
  - Easy to compute for small  $d$  and simple problems
  - For DL,  $H_f$  maybe too large and expensive to compute via backprop
- What is the best  $\epsilon$ ?



# Newton's method

- Reviewing the Taylor expansion of  $f$  there is no need to stop after the first term. In fact, we can write it as

$$f(\mathbf{x} + \epsilon) = f(\mathbf{x}) + \epsilon^\top \nabla f(\mathbf{x}) + \frac{1}{2} \epsilon^\top \nabla \nabla^\top f(\mathbf{x}) \epsilon + O(\|\epsilon\|^3)$$

- We define  $H_f := \Delta \Delta^\top f(\mathbf{x})$  to be the Hessian of  $f$ 
  - Easy to compute for small  $d$  and simple problems
  - For DL,  $H_f$  maybe too large and expensive to compute via backprop
- What is the best  $\epsilon$ ?
- Consider  $f(x) = \frac{1}{2}x^2$ , what is the gradient descent rule using Newton's method?

# Newton's method

- Reviewing the Taylor expansion of  $f$  there is no need to stop after the first term. In fact, we can write it as

$$f(\mathbf{x} + \epsilon) = f(\mathbf{x}) + \epsilon^\top \nabla f(\mathbf{x}) + \frac{1}{2} \epsilon^\top \nabla \nabla^\top f(\mathbf{x}) \epsilon + O(\|\epsilon\|^3)$$

- We define  $H_f := \Delta \Delta^\top f(\mathbf{x})$  to be the Hessian of  $f$ 
  - Easy to compute for small  $d$  and simple problems
  - For DL,  $H_f$  maybe too large and expensive to compute via backprop
- What is the best  $\epsilon$ ?
- Consider  $f(x) = \frac{1}{2}x^2$ , what is the gradient descent rule using Newton's method?
- Preconditioning: computing the inverse of Hessian is expensive. So only use the diagonal entries of Hessian

# Stochastic Gradient Descent (SGD)

- In DL, the objective function is often the average of losses per example in the training data set

$$f(\mathbf{x}) = \frac{1}{n} \sum_{i=1}^n f_i(\mathbf{x}).$$

# Stochastic Gradient Descent (SGD)

- In DL, the objective function is often the average of losses per example in the training data set

$$f(\mathbf{x}) = \frac{1}{n} \sum_{i=1}^n f_i(\mathbf{x}).$$

- The gradient of the objective function at  $\mathbf{x}$  is computed as  $\nabla f(\mathbf{x}) = \frac{1}{n} \sum_{i=1}^n \nabla f_i(\mathbf{x})$ .

# Stochastic Gradient Descent (SGD)

- In DL, the objective function is often the average of losses per example in the training data set

$$f(\mathbf{x}) = \frac{1}{n} \sum_{i=1}^n f_i(\mathbf{x}).$$

- The gradient of the objective function at  $\mathbf{x}$  is computed as  $\nabla f(\mathbf{x}) = \frac{1}{n} \sum_{i=1}^n \nabla f_i(\mathbf{x})$ .
- Computation cost of each update is  $\mathcal{O}(n)$

# Stochastic Gradient Descent (SGD)

- In DL, the objective function is often the average of losses per example in the training data set

$$f(\mathbf{x}) = \frac{1}{n} \sum_{i=1}^n f_i(\mathbf{x}).$$

- The gradient of the objective function at  $\mathbf{x}$  is computed as  $\nabla f(\mathbf{x}) = \frac{1}{n} \sum_{i=1}^n \nabla f_i(\mathbf{x})$ .
- Computation cost of each update is  $\mathcal{O}(n)$
- SGD reduces the computational cost at each iteration
  - At each iteration of SGD, we uniformly sample an index  $i \in \{1, \dots, n\}$  for data instances at random
  - Compute the gradient  $\nabla f_i(\mathbf{x})$  to update  $\mathbf{x}$

# Stochastic Gradient Descent (SGD)

- What is the computation cost per update?


# Stochastic Gradient Descent (SGD)

- What is the computation cost per update?
- Is the stochastic gradient  $\nabla f_i(\mathbf{x})$  unbiased estimate of  $\nabla f(\mathbf{x})$  ?



# Stochastic Gradient Descent (SGD)

- What is the computation cost per update?
- Is the stochastic gradient  $\nabla f_i(\mathbf{x})$  unbiased estimate of  $\nabla f(\mathbf{x})$  ?
- The trajectory of SGD looks more noisy than GD



# Mini-batch Stochastic Gradient Descent

- We can perform random uniform sampling for each iteration to form a minibatch and then use this minibatch to compute the gradient

# Mini-batch Stochastic Gradient Descent

- We can perform random uniform sampling for each iteration to form a minibatch and then use this minibatch to compute the gradient
- In each subsequent time step  $t > 0$ , mini-batch SGD uses random uniform sampling to get a mini-batch  $\mathcal{B}_t$  made up of example indices from the training set

# Mini-batch Stochastic Gradient Descent

- We can perform random uniform sampling for each iteration to form a minibatch and then use this minibatch to compute the gradient
- In each subsequent time step  $t > 0$ , mini-batch SGD uses random uniform sampling to get a mini-batch  $\mathcal{B}_t$  made up of example indices from the training set
- Then we compute the gradient  $\mathbf{g}_t$  of the objective function at  $\mathbf{x}_{t-1}$  with mini-batch  $\mathcal{B}_t$  at time step  $t$ :

$$\mathbf{g}_t \leftarrow \nabla f_{\mathcal{B}_t}(\mathbf{x}_{t-1}) = \frac{1}{|\mathcal{B}|} \sum_{i \in \mathcal{B}_t} \nabla f_i(\mathbf{x}_{t-1})$$

# Mini-batch Stochastic Gradient Descent

- We can perform random uniform sampling for each iteration to form a minibatch and then use this minibatch to compute the gradient
- In each subsequent time step  $t > 0$ , mini-batch SGD uses random uniform sampling to get a mini-batch  $\mathcal{B}_t$  made up of example indices from the training set
- Then we compute the gradient  $\mathbf{g}_t$  of the objective function at  $\mathbf{x}_{t-1}$  with mini-batch  $\mathcal{B}_t$  at time step  $t$ :

$$\mathbf{g}_t \leftarrow \nabla f_{\mathcal{B}_t}(\mathbf{x}_{t-1}) = \frac{1}{|\mathcal{B}|} \sum_{i \in \mathcal{B}_t} \nabla f_i(\mathbf{x}_{t-1})$$

- Given the learning rate  $\eta_t$ , the iteration of the mini-batch SGD is:

$$\mathbf{x}_t \leftarrow \mathbf{x}_{t-1} - \eta_t \mathbf{g}_t.$$

# Mini-batch Stochastic Gradient Descent

- We can perform random uniform sampling for each iteration to form a minibatch and then use this minibatch to compute the gradient
- In each subsequent time step  $t > 0$ , mini-batch SGD uses random uniform sampling to get a mini-batch  $\mathcal{B}_t$  made up of example indices from the training set
- Then we compute the gradient  $\mathbf{g}_t$  of the objective function at  $\mathbf{x}_{t-1}$  with mini-batch  $\mathcal{B}_t$  at time step  $t$ :

$$\mathbf{g}_t \leftarrow \nabla f_{\mathcal{B}_t}(\mathbf{x}_{t-1}) = \frac{1}{|\mathcal{B}|} \sum_{i \in \mathcal{B}_t} \nabla f_i(\mathbf{x}_{t-1})$$

- Given the learning rate  $\eta_t$ , the iteration of the mini-batch SGD is:

$$\mathbf{x}_t \leftarrow \mathbf{x}_{t-1} - \eta_t \mathbf{g}_t.$$

- SGD can self-decay itself by using  $\eta_t = \eta t^\alpha$  (usually  $\alpha = -1$  or  $\alpha = -0.5$ ),  $\eta_t = \eta \alpha^t$  (e.g.  $\alpha = 0.95$ )

# Momentum

- Let's start with an example  $f(\mathbf{x}) = 0.1 x_1^2 + 2x_2^2$


# Momentum

- Let's start with an example  $f(\mathbf{x}) = 0.1 x_1^2 + 2x_2^2$
- Let's implement GD on this objective function and demonstrate the iterative trajectory of the updates with learning rate 0.4




# Momentum

- Let's start with an example  $f(\mathbf{x}) = 0.1 x_1^2 + 2x_2^2$
- Let's implement GD on this objective function and demonstrate the iterative trajectory of the updates with learning rate 0.4
  - The slope of the objective function has a larger value in the vertical direction
  - Therefore, the variable will move more in Vertical direction




# Momentum

- Let's start with an example  $f(\mathbf{x}) = 0.1 x_1^2 + 2x_2^2$
- Let's implement GD on this objective function and demonstrate the iterative trajectory of the updates with learning rate 0.4
  - The slope of the objective function has a larger value in the vertical direction
  - Therefore, the variable will move more in Vertical direction
- Small learning rate will cause the variable move slower toward the optimal solution



# Momentum

- Let's start with an example  $f(\mathbf{x}) = 0.1 x_1^2 + 2x_2^2$
- Let's implement GD on this objective function and demonstrate the iterative trajectory of the updates with learning rate 0.4
  - The slope of the objective function has a larger value in the vertical direction
  - Therefore, the variable will move more in Vertical direction
- Small learning rate will cause the variable move slower toward the optimal solution
- Large learning rate will make the variable overshoot in vertical direction



# Momentum

- Momentum method was proposed to address this problem

# Momentum

- Momentum method was proposed to address this problem
- At time step  $t > 0$ , momentum modifies the steps at each iteration as follows

$$\mathbf{v}_t \leftarrow \gamma \mathbf{v}_{t-1} + \eta_t \mathbf{g}_t,$$

$$\mathbf{x}_t \leftarrow \mathbf{x}_{t-1} - \mathbf{v}_t,$$

# Momentum

- Momentum method was proposed to address this problem
- At time step  $t > 0$ , momentum modifies the steps at each iteration as follows

$$\mathbf{v}_t \leftarrow \gamma \mathbf{v}_{t-1} + \eta_t \mathbf{g}_t,$$

$$\mathbf{x}_t \leftarrow \mathbf{x}_{t-1} - \mathbf{v}_t,$$

- When we implement this on the previous example we see that the trajectory is smoother in vertical direction and approaches the optimal solution faster in the horizontal direction.


# Momentum

- Momentum method was proposed to address this problem
- At time step  $t > 0$ , momentum modifies the steps at each iteration as follows

$$\mathbf{v}_t \leftarrow \gamma \mathbf{v}_{t-1} + \eta_t \mathbf{g}_t,$$

$$\mathbf{x}_t \leftarrow \mathbf{x}_{t-1} - \mathbf{v}_t,$$

- When we implement this on the previous example we see that the trajectory is smoother in vertical direction and approaches the optimal solution faster in the horizontal direction.



# Momentum

- To understand momentum, we can expanding the velocity variable  $\mathbf{v}_t$



# Momentum

- To understand momentum, we can expanding the velocity variable  $\mathbf{v}_t$  over time:

$$\begin{aligned}\mathbf{v}_t &= \eta_t \mathbf{g}_t + \gamma \mathbf{v}_{t-1}, \\ &= \eta_t \mathbf{g}_t + \gamma \eta_{t-1} \mathbf{g}_{t-1} + \gamma \mathbf{v}_{t-1}, \\ &\dots \\ &= \eta_t \mathbf{g}_t + \gamma \eta_{t-1} \mathbf{g}_{t-1} + \dots + \gamma^{t-1} \eta_1 \mathbf{g}_1.\end{aligned}$$

# Momentum

- To understand momentum, we can expanding the velocity variable  $\mathbf{v}_t$  over time:

$$\begin{aligned}\mathbf{v}_t &= \eta_t \mathbf{g}_t + \gamma \mathbf{v}_{t-1}, \\ &= \eta_t \mathbf{g}_t + \gamma \eta_{t-1} \mathbf{g}_{t-1} + \gamma \mathbf{v}_{t-1}, \\ &\dots \\ &= \eta_t \mathbf{g}_t + \gamma \eta_{t-1} \mathbf{g}_{t-1} + \dots + \gamma^{t-1} \eta_1 \mathbf{g}_1.\end{aligned}$$

- $\mathbf{v}_t$  is a weighted sum over all past gradients multiplied by the according learning rate

# Momentum

- To understand momentum, we can expanding the velocity variable  $\mathbf{v}_t$  over time:

$$\begin{aligned}\mathbf{v}_t &= \eta_t \mathbf{g}_t + \gamma \mathbf{v}_{t-1}, \\ &= \eta_t \mathbf{g}_t + \gamma \eta_{t-1} \mathbf{g}_{t-1} + \gamma \mathbf{v}_{t-1}, \\ &\dots \\ &= \eta_t \mathbf{g}_t + \gamma \eta_{t-1} \mathbf{g}_{t-1} + \dots + \gamma^{t-1} \eta_1 \mathbf{g}_1.\end{aligned}$$


- $\mathbf{v}_t$  is a weighted sum over all past gradients multiplied by the according learning rate
- What happens when  $\gamma$  is too large or small?

# Momentum

- To understand momentum, we can expand the velocity variable  $\mathbf{v}_t$  over time:

$$\begin{aligned}\mathbf{v}_t &= \eta_t \mathbf{g}_t + \gamma \mathbf{v}_{t-1}, \\ &= \eta_t \mathbf{g}_t + \gamma \eta_{t-1} \mathbf{g}_{t-1} + \gamma \mathbf{v}_{t-1}, \\ &\dots \\ &= \eta_t \mathbf{g}_t + \gamma \eta_{t-1} \mathbf{g}_{t-1} + \dots + \gamma^{t-1} \eta_1 \mathbf{g}_1.\end{aligned}$$

- $\mathbf{v}_t$  is a weighted sum over all past gradients multiplied by the according learning rate
- What happens when  $\gamma$  is too large or small?



# Adagrad

- So far we showed that each variable uses the same learning rate

# Adagrad

- So far we showed that each variable uses the same learning rate
- For example, if we assume that the objective function is  $f$  and the independent variable is a two-dimensional vector  $[x_1, x_2]$ , the update for each element is  $x_1 \leftarrow x_1 - \eta \frac{\partial f}{\partial x_1}$      $x_2 \leftarrow x_2 - \eta \frac{\partial f}{\partial x_2}$

# Adagrad

- So far we showed that each variable uses the same learning rate
- For example, if we assume that the objective function is  $f$  and the independent variable is a two-dimensional vector  $[x_1, x_2]$ , the update for each element is  $x_1 \leftarrow x_1 - \eta \frac{\partial f}{\partial x_1}$      $x_2 \leftarrow x_2 - \eta \frac{\partial f}{\partial x_2}$
- We saw that this would be a problem when there is a big difference between  $x_1$  and  $x_2$ .

# Adagrad

- So far we showed that each variable uses the same learning rate
- For example, if we assume that the objective function is  $f$  and the independent variable is a two-dimensional vector  $[x_1, x_2]$ , the update for each element is  $x_1 \leftarrow x_1 - \eta \frac{\partial f}{\partial x_1}$      $x_2 \leftarrow x_2 - \eta \frac{\partial f}{\partial x_2}$
- We saw that this would be a problem when there is a big difference between  $x_1$  and  $x_2$ .
- The momentum relies on the exponentially weighted moving average to make the direction of the independent variable more consistent



# Adagrad

- So far we showed that each variable uses the same learning rate
- For example, if we assume that the objective function is  $f$  and the independent variable is a two-dimensional vector  $[x_1, x_2]$ , the update for each element is  $x_1 \leftarrow x_1 - \eta \frac{\partial f}{\partial x_1}$      $x_2 \leftarrow x_2 - \eta \frac{\partial f}{\partial x_2}$
- We saw that this would be a problem when there is a big difference between  $x_1$  and  $x_2$ .
- The momentum relies on the exponentially weighted moving average to make the direction of the independent variable more consistent
- Now, we introduce Adagrad that adjusts the learning rate according to the gradient value of the independent variable in each direction

# Adagrad

- Adagrad uses the cumulative variable  $\mathbf{s}_t$  obtained from a square by element operation on the mini-batch stochastic gradient  $\mathbf{g}_t$ .

# Adagrad

- Adagrad uses the cumulative variable  $\mathbf{s}_t$  obtained from a square by element operation on the mini-batch stochastic gradient  $\mathbf{g}_t$ .
- At time step 0,  $\mathbf{s}_0 = \mathbf{0}$ .

# Adagrad

- Adagrad uses the cumulative variable  $\mathbf{s}_t$  obtained from a square by element operation on the mini-batch stochastic gradient  $\mathbf{g}_t$ .
- At time step 0,  $\mathbf{s}_0 = 0$ .
- At time step  $t$ , we first sum the results of the square by element operation for the mini-batch gradient  $\mathbf{g}_t$  to get the variable  $\mathbf{s}_t$ :

$$\mathbf{s}_t \leftarrow \mathbf{s}_{t-1} + \mathbf{g}_t \odot \mathbf{g}_t,$$

# Adagrad

- Adagrad uses the cumulative variable  $\mathbf{s}_t$  obtained from a square by element operation on the mini-batch stochastic gradient  $\mathbf{g}_t$ .
- At time step 0,  $\mathbf{s}_0 = 0$ .
- At time step  $t$ , we first sum the results of the square by element operation for the mini-batch gradient  $\mathbf{g}_t$  to get the variable  $\mathbf{s}_t$ :

$$\mathbf{s}_t \leftarrow \mathbf{s}_{t-1} + \mathbf{g}_t \odot \mathbf{g}_t,$$

- Next, we readjust the learning rate of each element in the independent variable of the objective function using element operations

$$\mathbf{x}_t \leftarrow \mathbf{x}_{t-1} - \frac{\eta}{\sqrt{\mathbf{s}_t} + \epsilon} \odot \mathbf{g}_t,$$

# Adagrad

- if an element in the independent variable of the objective function has a constant and large partial derivative, the learning rate of this element will drop faster.

# Adagrad

- if an element in the independent variable of the objective function has a constant and large partial derivative, the learning rate of this element will drop faster.
- On the contrary, if the partial derivative of such an element remains small, then its learning rate will decline more slowly.



# Adagrad

- if an element in the independent variable of the objective function has a constant and large partial derivative, the learning rate of this element will drop faster.
- On the contrary, if the partial derivative of such an element remains small, then its learning rate will decline more slowly.
- when the learning rate declines very fast during early iteration, yet the current solution is still not desirable, Adagrad might have difficulty finding a useful solution because the learning rate will be too small at later stages of iteration.



# Adagrad

- if an element in the independent variable of the objective function has a constant and large partial derivative, the learning rate of this element will drop faster.
- On the contrary, if the partial derivative of such an element remains small, then its learning rate will decline more slowly.
- when the learning rate declines very fast during early iteration, yet the current solution is still not desirable, Adagrad might have difficulty finding a useful solution because the learning rate will be too small at later stages of iteration.
- If we use Adagrad for the example  $f(\mathbf{x}) = 0.1 x_1^2 + 2x_2^2$ , we get



Which one uses a larger learning rate?

# RMSProp

- RMSProp made a small modification to Adagrad to prevent learning rate from declining very fast during early iteration

# RMSProp

- RMSProp made a small modification to Adagrad to prevent learning rate from declining very fast during early iteration
- RMSProp uses the exponentially weighted moving average on the square by element results of these gradients.

# RMSProp

- RMSProp made a small modification to Adagrad to prevent learning rate from declining very fast during early iteration
- RMSProp uses the exponentially weighted moving average on the square by element results of these gradients.
- Specifically, given the hyperparameter  $0 \leq \gamma < 1$ , RMSProp is computed at time step  $t > 0$ :

$$\mathbf{s}_t \leftarrow \gamma \mathbf{s}_{t-1} + (1 - \gamma) \mathbf{g}_t \odot \mathbf{g}_t.$$

# RMSProp

- RMSProp made a small modification to Adagrad to prevent learning rate from declining very fast during early iteration
- RMSProp uses the exponentially weighted moving average on the square by element results of these gradients.
- Specifically, given the hyperparameter  $0 \leq \gamma < 1$ , RMSProp is computed at time step  $t > 0$ :

$$\mathbf{s}_t \leftarrow \gamma \mathbf{s}_{t-1} + (1 - \gamma) \mathbf{g}_t \odot \mathbf{g}_t.$$

- Like Adagrad, RMSProp readjust the learning rate of each element in the independent variable of the object function as:

$$\mathbf{x}_t \leftarrow \mathbf{x}_{t-1} - \frac{\eta}{\sqrt{\mathbf{s}_t + \epsilon}} \odot \mathbf{g}_t,$$

# RMSProp

- If we expand the definition of  $\mathbf{s}_t$ , we see that:


$$\begin{aligned}\mathbf{s}_t &= (1 - \gamma)\mathbf{g}_t \odot \mathbf{g}_t + \gamma\mathbf{s}_{t-1} \\ &= (1 - \gamma)(\mathbf{g}_t \odot \mathbf{g}_t + \gamma\mathbf{g}_{t-1} \odot \mathbf{g}_{t-1}) + \gamma^2\mathbf{s}_{t-2} \\ &\dots \\ &= (1 - \gamma)(\mathbf{g}_t \odot \mathbf{g}_t + \gamma\mathbf{g}_{t-1} \odot \mathbf{g}_{t-1} + \dots + \gamma^{t-1}\mathbf{g}_1 \odot \mathbf{g}_1) .\end{aligned}$$

# RMSProp

- If we expand the definition of  $\mathbf{s}_t$ , we see that:

$$\begin{aligned}\mathbf{s}_t &= (1 - \gamma)\mathbf{g}_t \odot \mathbf{g}_t + \gamma\mathbf{s}_{t-1} \\ &= (1 - \gamma)(\mathbf{g}_t \odot \mathbf{g}_t + \gamma\mathbf{g}_{t-1} \odot \mathbf{g}_{t-1}) + \gamma^2\mathbf{s}_{t-2} \\ &\dots \\ &= (1 - \gamma)(\mathbf{g}_t \odot \mathbf{g}_t + \gamma\mathbf{g}_{t-1} \odot \mathbf{g}_{t-1} + \dots + \gamma^{t-1}\mathbf{g}_1 \odot \mathbf{g}_1) .\end{aligned}$$

- We visualize these weights in the past 40 time steps for various  $\gamma$ :



# Adadelta

- In addition to RMSProp, Adadelta is another common optimization algorithm that helps improve the chances of finding useful solutions at later stages of iteration.



# Adadelta

- In addition to RMSProp, Adadelta is another common optimization algorithm that helps improve the chances of finding useful solutions at later stages of iteration.
- The interesting thing is that there is no learning rate hyperparameter in the Adadelta algorithm.

# Adadelta

- In addition to RMSProp, Adadelta is another common optimization algorithm that helps improve the chances of finding useful solutions at later stages of iteration.
- The interesting thing is that there is no learning rate hyperparameter in the Adadelta algorithm.
- Given the hyperparameter  $0 \leq \rho < 1$ , we compute  $s_t \leftarrow \rho s_{t-1} + (1 - \rho) \mathbf{g}_t \odot \mathbf{g}_t$ .

# Adadelta

- In addition to RMSProp, Adadelta is another common optimization algorithm that helps improve the chances of finding useful solutions at later stages of iteration.
- The interesting thing is that there is no learning rate hyperparameter in the Adadelta algorithm.
- Given the hyperparameter  $0 \leq \rho < 1$ , we compute  $s_t \leftarrow \rho s_{t-1} + (1 - \rho) g_t \odot g_t$ .
- Unlike RMSprop, Adadelta maintains an additional state variable  $\Delta \mathbf{x}_t$  to compute the variation of the independent variable

$$g'_t \leftarrow \sqrt{\frac{\Delta \mathbf{x}_{t-1} + \epsilon}{s_t + \epsilon}} \odot g_t,$$

# Adadelta

- In addition to RMSProp, Adadelta is another common optimization algorithm that helps improve the chances of finding useful solutions at later stages of iteration.
- The interesting thing is that there is no learning rate hyperparameter in the Adadelta algorithm.
- Given the hyperparameter  $0 \leq \rho < 1$ , we compute  $s_t \leftarrow \rho s_{t-1} + (1 - \rho) g_t \odot g_t$ .
- Unlike RMSprop, Adadelta maintains an additional state variable  $\Delta \mathbf{x}_t$  to compute the variation of the independent variable

$$g'_t \leftarrow \sqrt{\frac{\Delta \mathbf{x}_{t-1} + \epsilon}{s_t + \epsilon}} \odot g_t,$$

- Next we update the independent variable  $\mathbf{x}_t \leftarrow \mathbf{x}_{t-1} - g'_t$ .

# Adadelta

- In addition to RMSProp, Adadelta is another common optimization algorithm that helps improve the chances of finding useful solutions at later stages of iteration.
- The interesting thing is that there is no learning rate hyperparameter in the Adadelta algorithm.
- Given the hyperparameter  $0 \leq \rho < 1$ , we compute  $s_t \leftarrow \rho s_{t-1} + (1 - \rho) g_t \odot g_t$ .
- Unlike RMSprop, Adadelta maintains an additional state variable  $\Delta \mathbf{x}_t$  to compute the variation of the independent variable

$$g'_t \leftarrow \sqrt{\frac{\Delta \mathbf{x}_{t-1} + \epsilon}{s_t + \epsilon}} \odot g_t,$$

- Next we update the independent variable  $\mathbf{x}_t \leftarrow \mathbf{x}_{t-1} - g'_t$ .
- Finally, we use  $\Delta \mathbf{x}_t$  to record EWMA on the squares of elements of  $\mathbf{g}'$

$$\Delta \mathbf{x}_t \leftarrow \rho \Delta \mathbf{x}_{t-1} + (1 - \rho) g'_t \odot g'_t.$$

# Adam

- Created on the basis of RMSProp, Adam also uses EWMA on the mini-batch stochastic gradient

# Adam

- Created on the basis of RMSProp, Adam also uses EWMA on the mini-batch stochastic gradient
- Adam uses a momentum variable  $\mathbf{v}_t$  and variable  $\mathbf{s}_t$ , which is an EWMA on the squares of elements in the mini-batch SGD from RMSProp

$$\mathbf{v}_t \leftarrow \beta_1 \mathbf{v}_{t-1} + (1 - \beta_1) \mathbf{g}_t. \quad \mathbf{s}_t \leftarrow \beta_2 \mathbf{s}_{t-1} + (1 - \beta_2) \mathbf{g}_t \odot \mathbf{g}_t.$$

# Adam

- Created on the basis of RMSProp, Adam also uses EWMA on the mini-batch stochastic gradient
- Adam uses a momentum variable  $\mathbf{v}_t$  and variable  $\mathbf{s}_t$ , which is an EWMA on the squares of elements in the mini-batch SGD from RMSProp

$$\mathbf{v}_t \leftarrow \beta_1 \mathbf{v}_{t-1} + (1 - \beta_1) \mathbf{g}_t. \quad \mathbf{s}_t \leftarrow \beta_2 \mathbf{s}_{t-1} + (1 - \beta_2) \mathbf{g}_t \odot \mathbf{g}_t.$$

- Notice that  $\mathbf{v}_t = \underbrace{(1 - \beta_1) \sum_{i=1}^t \beta_1^{t-i}}_{1 - \beta_1^t} \mathbf{g}_i$  What happens when  $t$  is small?



# Adam

- Created on the basis of RMSProp, Adam also uses EWMA on the mini-batch stochastic gradient
- Adam uses a momentum variable  $\mathbf{v}_t$  and variable  $\mathbf{s}_t$ , which is an EWMA on the squares of elements in the mini-batch SGD from RMSProp

$$\mathbf{v}_t \leftarrow \beta_1 \mathbf{v}_{t-1} + (1 - \beta_1) \mathbf{g}_t. \quad \mathbf{s}_t \leftarrow \beta_2 \mathbf{s}_{t-1} + (1 - \beta_2) \mathbf{g}_t \odot \mathbf{g}_t.$$

- Notice that  $\mathbf{v}_t = \underbrace{(1 - \beta_1) \sum_{i=1}^t \beta_1^{t-i}}_{1 - \beta_1^t} \mathbf{g}_i$  What happens when  $t$  is small?

$$1 - \beta_1^t$$

- When  $t$  is small, the sum of the mini-batch stochastic gradient weights from each previous time step will be small. To eliminate this effect, we perform bias correction:

$$\hat{\mathbf{v}}_t \leftarrow \frac{\mathbf{v}_t}{1 - \beta_1^t}, \quad \hat{\mathbf{s}}_t \leftarrow \frac{\mathbf{s}_t}{1 - \beta_2^t}.$$

# Adam

- Next, the Adam algorithm will use the bias-corrected variables to re-adjust the learning rate of each element.

$$\mathbf{g}'_t \leftarrow \frac{\eta \hat{\mathbf{v}}_t}{\sqrt{\hat{\mathbf{s}}_t} + \epsilon},$$

# Adam

- Next, the Adam algorithm will use the bias-corrected variables to re-adjust the learning rate of each element.

$$\mathbf{g}'_t \leftarrow \frac{\eta \hat{\mathbf{v}}_t}{\sqrt{\hat{\mathbf{s}}_t} + \epsilon},$$

- Finally, we update the independent variable as

$$\mathbf{x}_t \leftarrow \mathbf{x}_{t-1} - \mathbf{g}'_t.$$

# Convolutional Neural Networks

- So far, we simply threw away the spatial structure in images by flattening each image into a 1D vector, and fed it into a fully-connected network

# Convolutional Neural Networks

- So far, we simply threw away the spatial structure in images by flattening each image into a 1D vector, and fed it into a fully-connected network
- Ideally, we would find a way to leverage our prior knowledge that nearby pixels are more related to each other.

# Convolutional Neural Networks

- So far, we simply threw away the spatial structure in images by flattening each image into a 1D vector, and fed it into a fully-connected network
- Ideally, we would find a way to leverage our prior knowledge that nearby pixels are more related to each other.
- Now, we introduce convolutional neural networks (CNNs), a powerful family of neural networks that were designed for precisely this purpose.

# Convolutional Neural Networks

- So far, we simply threw away the spatial structure in images by flattening each image into a 1D vector, and fed it into a fully-connected network
- Ideally, we would find a way to leverage our prior knowledge that nearby pixels are more related to each other.
- Now, we introduce convolutional neural networks (CNNs), a powerful family of neural networks that were designed for precisely this purpose.
- In addition to their strong predictive performance,
  - convolutional neural networks tend to be computationally efficient,
  - both because they tend to require fewer parameters than dense architectures
  - also because convolutions are easy to parallelize across GPU cores

# From Dense Layers to Convolutions

- A few key principles for building neural networks for computer vision:
  1. **Translation Invariance:** Our vision systems should, in some sense, respond similarly to the same object regardless of where it appears in the image



# From Dense Layers to Convolutions

- A few key principles for building neural networks for computer vision:
  1. **Translation Invariance:** Our vision systems should, in some sense, respond similarly to the same object regardless of where it appears in the image
  2. **Locality:** Our vision systems should, in some sense, focus on local regions, without regard for what else is happening in the image at greater distances.

# Constraining the MLP

- Let's consider how MLP would look like with  $h \times w$  images as inputs, and hidden representations as  $h \times w$  matrices.

# Constraining the MLP

- Let's consider how MLP would look like with  $h \times w$  images as inputs, and hidden representations as  $h \times w$  matrices.
- To have each of the  $hw$  hidden nodes receive input from each of the  $hw$  inputs, we would switch from using weight matrices to representing our parameters as four-dimensional weight tensors.

# Constraining the MLP

- Let's consider how MLP would look like with  $h \times w$  images as inputs, and hidden representations as  $h \times w$  matrices.
- To have each of the  $hw$  hidden nodes receive input from each of the  $hw$  inputs, we would switch from using weight matrices to representing our parameters as four-dimensional weight tensors.
- We could formally express this dense layer as follows:

$$h[i, j] = u[i, j] + \sum_{k, l} W[i, j, k, l] \cdot x[k, l] = u[i, j] + \sum_{a, b} V[i, j, a, b] \cdot x[i + a, j + b]$$

# Constraining the MLP

- Let's consider how MLP would look like with  $h \times w$  images as inputs, and hidden representations as  $h \times w$  matrices.
- To have each of the  $hw$  hidden nodes receive input from each of the  $hw$  inputs, we would switch from using weight matrices to representing our parameters as four-dimensional weight tensors.
- We could formally express this dense layer as follows:

$$h[i, j] = u[i, j] + \sum_{k, l} W[i, j, k, l] \cdot x[k, l] = u[i, j] + \sum_{a, b} V[i, j, a, b] \cdot x[i + a, j + b]$$

- For any given location  $(i, j)$  in the hidden layer  $h[i, j]$ , we compute its value by summing over pixels in  $x$ , centered around  $(i, j)$  and weighted by  $V[i, j, a, b]$ .

# Constraining the MLP

- Let's invoke the first principle- **translation invariance**.

# Constraining the MLP

- Let's invoke the first principle- **translation invariance**.
- This implies that a shift in the inputs  $x$  should simply lead to a shift in the activations  $h$ .

# Constraining the MLP

- Let's invoke the first principle- **translation invariance**.
- This implies that a shift in the inputs  $x$  should simply lead to a shift in the activations  $h$ .
- This is only possible if  $V$  and  $u$  don't actually depend on  $(i, j)$ .



# Constraining the MLP

- Let's invoke the first principle- **translation invariance**.
- This implies that a shift in the inputs  $x$  should simply lead to a shift in the activations  $h$ .
- This is only possible if  $V$  and  $u$  don't actually depend on  $(i, j)$ .
- In other words, we have  $V[i, j, a, b] = V[a, b]$  and  $u$  is a constant.

# Constraining the MLP

- Let's invoke the first principle- **translation invariance**.
- This implies that a shift in the inputs  $x$  should simply lead to a shift in the activations  $h$ .
- This is only possible if  $V$  and  $u$  don't actually depend on  $(i, j)$ .
- In other words, we have  $V[i, j, a, b] = V[a, b]$  and  $u$  is a constant.
- As a result we can simplify the definition for  $h$ :

$$h[i, j] = u + \sum_{a, b} V[a, b] \cdot x[i + a, j + b]$$

# Constraining the MLP

- Let's invoke the first principle- **translation invariance**.
- This implies that a shift in the inputs  $x$  should simply lead to a shift in the activations  $h$ .
- This is only possible if  $V$  and  $u$  don't actually depend on  $(i, j)$ .
- In other words, we have  $V[i, j, a, b] = V[a, b]$  and  $u$  is a constant.
- As a result we can simplify the definition for  $h$ :

$$h[i, j] = u + \sum_{a, b} V[a, b] \cdot x[i + a, j + b]$$

- This is a convolution! We also reduced the number of parameters.

# Constraining the MLP

- Now let's invoke the second principle – **locality**.

# Constraining the MLP

- Now let's invoke the second principle – **locality**.
- We believe we should not have to look very far away from  $(i, j)$  in order to glean relevant information to assess what is going on at  $h[i, j]$ .

# Constraining the MLP

- Now let's invoke the second principle – **locality**.
- We believe we should not have to look very far away from  $(i, j)$  in order to glean relevant information to assess what is going on at  $h[i, j]$ .
- This means that outside some range  $|a|, |b| > \Delta$ , we should set  $V[a, b] = 0$ .

# Constraining the MLP

- Now let's invoke the second principle – **locality**.
- We believe we should not have to look very far away from  $(i, j)$  in order to glean relevant information to assess what is going on at  $h[i, j]$ .
- This means that outside some range  $|a|, |b| > \Delta$ , we should set  $V[a, b] = 0$ .
- Equivalently, we can write  $h[i, j]$  as

$$h[i, j] = u + \sum_{a=-\Delta}^{\Delta} \sum_{b=-\Delta}^{\Delta} V[a, b] \cdot x[i + a, j + b]$$

# Constraining the MLP

- Now let's invoke the second principle – **locality**.
- We believe we should not have to look very far away from  $(i, j)$  in order to glean relevant information to assess what is going on at  $h[i, j]$ .
- This means that outside some range  $|a|, |b| > \Delta$ , we should set  $V[a, b] = 0$ .
- Equivalently, we can write  $h[i, j]$  as

$$h[i, j] = u + \sum_{a=-\Delta}^{\Delta} \sum_{b=-\Delta}^{\Delta} V[a, b] \cdot x[i + a, j + b]$$

- This, in a nutshell is the convolutional layer.



# Convolutions

- In mathematics, the convolution between two functions is defined as:

$$[f \circledast g](x) = \int_{\mathbb{R}^d} f(z)g(x - z)dz$$

# Convolutions

- In mathematics, the convolution between two functions is defined as:

$$[f \circledast g](x) = \int_{\mathbb{R}^d} f(z)g(x - z)dz$$

- For discrete variables, the integral turns into a sum

$$[f \circledast g](i) = \sum f(a)g(i - a)$$

# Convolutions

- In mathematics, the convolution between two functions is defined as:

$$[f \circledast g](x) = \int_{\mathbb{R}^d} f(z)g(x - z)dz$$

- For discrete variables, the integral turns into a sum

$$[f \circledast g](i) = \sum_a f(a)g(i - a)$$

- For two dimensional arrays, we have a corresponding sum with indices  $(i, j)$  for  $f$  and  $(i - a, j - b)$  for  $g$  respectively.

# Convolutions

- In mathematics, the convolution between two functions is defined as:

$$[f \circledast g](x) = \int_{\mathbb{R}^d} f(z)g(x - z)dz$$

- For discrete variables, the integral turns into a sum

$$[f \circledast g](i) = \sum_a f(a)g(i - a)$$

- For two dimensional arrays, we have a corresponding sum with indices  $(i, j)$  for  $f$  and  $(i - a, j - b)$  for  $g$  respectively.

- This looks similar to the definition  $h[i, j] = u + \sum_{a, b} V[a, b] \cdot x[i + a, j + b]$

# Convolutions

- In mathematics, the convolution between two functions is defined as:

$$[f \circledast g](x) = \int_{\mathbb{R}^d} f(z)g(x - z)dz$$

- For discrete variables, the integral turns into a sum

$$[f \circledast g](i) = \sum_a f(a)g(i - a)$$

- For two dimensional arrays, we have a corresponding sum with indices  $(i, j)$  for  $f$  and  $(i - a, j - b)$  for  $g$  respectively.

- This looks similar to the definition  $h[i, j] = u + \sum_{a, b} V[a, b] \cdot x[i + a, j + b]$

- In order to match the signs, we can use  $\tilde{V} = V[-a, -b]$  to obtain  $h = x \circledast \tilde{V}$

# Convolutions

- In mathematics, the convolution between two functions is defined as:

$$[f \circledast g](x) = \int_{\mathbb{R}^d} f(z)g(x - z)dz$$

- For discrete variables, the integral turns into a sum

$$[f \circledast g](i) = \sum_a f(a)g(i - a)$$

- For two dimensional arrays, we have a corresponding sum with indices  $(i, j)$  for  $f$  and  $(i - a, j - b)$  for  $g$  respectively.

- This looks similar to the definition  $h[i, j] = u + \sum_{a, b} V[a, b] \cdot x[i + a, j + b]$

- In order to match the signs, we can use  $\tilde{V} = V[-a, -b]$  to obtain  $h = x \circledast \tilde{V}$

- Also note that the original definition is actually a *cross-correlation*.

# Convolutions for Images

- Let's see how convolutions work in practice.

Input                      Kernel                      Output

0	1	2
3	4	5
6	7	8

\*

0	1
2	3

=

19	25
37	43

# Convolutions for Images

- Let's see how convolutions work in practice.

Input                      Kernel                      Output

0	1	2
3	4	5
6	7	8

\*

0	1
2	3

=

19	25
37	43

- Start the notebook IN\_CLASS\_convolution



# Padding and Stride

- In general, assuming the input shape is  $(n_h, n_w)$  and the convolution kernel window shape is  $(k_h, k_w)$ , then the output shape will be:

# Padding and Stride

- In general, assuming the input shape is  $(n_h, n_w)$  and the convolution kernel window shape is  $(k_h, k_w)$ , then the output shape will be:

$$(n_h - k_h + 1) \times (n_w - k_w + 1).$$

# Padding and Stride

- In general, assuming the input shape is  $(n_h, n_w)$  and the convolution kernel window shape is  $(k_h, k_w)$ , then the output shape will be:

$$(n_h - k_h + 1) \times (n_w - k_w + 1).$$

- This might result in having much smaller images at the output of convolutional layer and we might lose any interesting information in the boundaries.

# Padding and Stride

- In general, assuming the input shape is  $(n_h, n_w)$  and the convolution kernel window shape is  $(k_h, k_w)$ , then the output shape will be:

$$(n_h - k_h + 1) \times (n_w - k_w + 1).$$

- This might result in having much smaller images at the output of convolutional layer and we might lose any interesting information in the boundaries.
  - Solution: **Padding**

# Padding and Stride

- In general, assuming the input shape is  $(n_h, n_w)$  and the convolution kernel window shape is  $(k_h, k_w)$ , then the output shape will be:

$$(n_h - k_h + 1) \times (n_w - k_w + 1).$$

- This might result in having much smaller images at the output of convolutional layer and we might lose any interesting information in the boundaries.
  - Solution: **Padding**
- In some cases, we want to reduce the resolution drastically

# Padding and Stride

- In general, assuming the input shape is  $(n_h, n_w)$  and the convolution kernel window shape is  $(k_h, k_w)$ , then the output shape will be:

$$(n_h - k_h + 1) \times (n_w - k_w + 1).$$

- This might result in having much smaller images at the output of convolutional layer and we might lose any interesting information in the boundaries.
  - Solution: **Padding**
- In some cases, we want to reduce the resolution drastically
  - Solution: **Strides**

# Padding

- Adding extra pixels of filler around the boundary of our input image, thus increasing the effective size of the image.

Input                      Kernel                      Output

0	0	0	0	0
0	0	1	2	0
0	3	4	5	0
0	6	7	8	0
0	0	0	0	0

\*

0	1
2	3

=

0	3	8	4
9	19	25	10
21	37	43	16
6	7	8	0

# Padding

- Adding extra pixels of filler around the boundary of our input image, thus increasing the effective size of the image.

Input                      Kernel                      Output

0	0	0	0	0
0	0	1	2	0
0	3	4	5	0
0	6	7	8	0
0	0	0	0	0

\*

0	1
2	3

=

0	3	8	4
9	19	25	10
21	37	43	16
6	7	8	0

- What is the size of the output after padding?



# Padding

- Adding extra pixels of filler around the boundary of our input image, thus increasing the effective size of the image.

Input                      Kernel                      Output

0	0	0	0	0
0	0	1	2	0
0	3	4	5	0
0	6	7	8	0
0	0	0	0	0

\*

0	1
2	3

=

0	3	8	4
9	19	25	10
21	37	43	16
6	7	8	0

- What is the size of the output after padding?
- What do you think is a good number of padding?

# Stride

- When computing the convolution, we start with the convolution window at the top-left corner of the input array, and then slide it over all locations both down and to the right.

# Stride

- When computing the convolution, we start with the convolution window at the top-left corner of the input array, and then slide it over all locations both down and to the right.
- In previous examples, we default to sliding one pixel at a time.

# Stride

- When computing the convolution, we start with the convolution window at the top-left corner of the input array, and then slide it over all locations both down and to the right.
- In previous examples, we default to sliding one pixel at a time.
- However, sometimes, we move our window more than one pixel at a time, skipping the intermediate locations.

# Stride

- When computing the convolution, we start with the convolution window at the top-left corner of the input array, and then slide it over all locations both down and to the right.
- In previous examples, we default to sliding one pixel at a time.
- However, sometimes, we move our window more than one pixel at a time, skipping the intermediate locations.
- We refer to the number of rows and columns traversed per slide as the *stride*.

Input                      Kernel                      Output

0	0	0	0	0
0	0	1	2	0
0	3	4	5	0
0	6	7	8	0
0	0	0	0	0

\*      

0	1
2	3

      =      

0	8
6	8

Example for strides with 3 and 2 for height and width, respectively.